

Document routing decides which extraction pipeline a file deserves.

Fig. 3a

A converter that ignores page geometry will interleave columns of text.

Fig. 3b

The probe reads the file header instead of trusting the extension.

Table 2

Layout is measured per horizontal band, not from the page as a whole.

$n = 41$

A scanned page carries no text layer and must be rendered as an image.

$P < .05$

Structure is lost when a layout-aware reader is replaced by a naive one.

et al.

Magic bytes are authoritative because extensions are frequently wrong.

12.5%

see Fig. 1

Every decision here is deterministic, offline, and free of model calls.

20 μm

RT

$\sim 3 \text{ eV}$

$\times 10^4$

The gutter between columns is found by merging text intervals per band.

Superscripts, italics and headings survive a layout-aware conversion.

A single reference page must not outvote multi-column body pages.

The correct pipeline for a two-column paper is the layout-aware one.