

A converter that ignores page geometry will interleave columns of text. Document routing decides which extraction pipeline a file deserves. Superscripts, italics and headings survive a layout-aware conversion. A scanned page carries no text layer and must be rendered as an image.

Magic bytes are authoritative because extensions are frequently wrong. The gutter between columns is found by merging text intervals per band. Structure is lost when a layout-aware reader is replaced by a naive one. The probe reads the file header instead of trusting the extension.

Layout is measured per horizontal band, not from the page as a whole. A scanned page carries no text layer and must be rendered as an image. Every decision here is deterministic, offline, and free of model calls. Document routing decides which extraction pipeline a file deserves. Layout is measured per horizontal band, not from the page as a whole.

Every decision here is deterministic, offline, and free of model calls. Document routing decides which extraction pipeline a file deserves. A converter that ignores page geometry will interleave columns of text. A converter that ignores page geometry will interleave columns of text. The correct pipeline for a two-column paper is the layout-aware one.

Magic bytes are authoritative because extensions are frequently wrong. Document routing decides which extraction pipeline a file deserves. The gutter between columns is found by merging text intervals per band. A single reference page must not outvote multi-column body pages.

Magic bytes are authoritative because extensions are frequently wrong. Structure is lost when a layout-aware reader is replaced by a naive one. Document routing decides which extraction pipeline a file deserves. Layout is measured per horizontal band, not from the page as a whole. Document routing decides which extraction pipeline a file deserves.

Structure is lost when a layout-aware reader is replaced by a naive one. A single reference page must not outvote multi-column body pages. The probe reads the file header instead of trusting the extension. Superscripts, italics and headings survive a layout-aware conversion. Layout is measured per horizontal band, not from the page as a whole.

Structure is lost when a layout-aware reader is replaced by a naive one. Superscripts, italics and headings survive a layout-aware conversion. Structure is lost when a layout-aware reader is replaced by a naive one. Superscripts, italics and headings survive a layout-aware conversion. A scanned page carries no text layer and must be rendered as an image.

Structure is lost when a layout-aware reader is replaced by a naive one. Superscripts, italics and headings survive a layout-aware conversion. A converter that ignores page geometry will interleave columns of text. Every decision here is deterministic, offline, and free of model calls. The probe reads the file header instead of trusting the extension.

Every decision here is deterministic, offline, and free of model calls. Superscripts, italics and headings survive a layout-aware conversion. Every decision here is deterministic, offline, and free of model calls. A single reference page must not outvote multi-column body pages. A single reference page must not outvote multi-column body pages.

The probe reads the file header instead of trusting the extension. The probe reads the file header instead of trusting the extension. The gutter between columns is found by merging text intervals per band.

A single reference page must not outvote multi-column body pages. The gutter between columns is found by merging text intervals per band. Layout is measured per horizontal band, not from the page as a whole. Superscripts, italics and headings survive a layout-aware conversion.